

RIDGE REGRESSION

In Python

Margarita Geleta – Oriol Narváz Serrano

Mathematical Optimization | Data Science

About the project

The aim of the project is to implement and solve in *Python* the constrained ridge regression model:

$$\min_{\mathbf{w}, \gamma} \frac{1}{2} (\mathbf{A}\mathbf{w} + \gamma\mathbf{e} - \mathbf{y})^T (\mathbf{A}\mathbf{w} + \gamma\mathbf{e} - \mathbf{y})$$

subject to $\|\mathbf{w}\|_2^2 \leq t$.

The data set we are going to use as an example is the *death rate* data set with 60 observations, where the *death rate* is represented as a function of 15 features as the *nitric oxide* pollution index, the *hydrocarbon* pollution index and so on.

In order to implement the constrained ridge regression model, we first need to find the formulae that define the objective function $f(\mathbf{w}, \gamma)$, the gradient of the objective function $\nabla f(\mathbf{w}, \gamma)$, the hessian of the objective function $\nabla^2 f(\mathbf{w}, \gamma)$, the inequality constraint $h(\mathbf{w})$, its gradient (the *Jacobian*) $\nabla h(\mathbf{w})$ and the hessian of the constraint $\nabla^2 h(\mathbf{w})$. So, we are going to work step by step upon these functions.

The objective function $f(\mathbf{w}, \gamma)$

It is clear that the objective function is the one we want to minimize, that is:

$$f(\mathbf{w}, \gamma) = \frac{1}{2} (\mathbf{A}\mathbf{w} + \gamma\mathbf{e} - \mathbf{y})^T (\mathbf{A}\mathbf{w} + \gamma\mathbf{e} - \mathbf{y})$$

To add some ease to the implementation, it is better to group \mathbf{w}, γ in just one variable $\boldsymbol{\omega}$ and thus, we add a column of ones in the matrix \mathbf{A} :

$$\mathbf{A}\mathbf{w} + \gamma\mathbf{e} = (\mathbf{A} \mid \mathbf{e}) \begin{pmatrix} \mathbf{w} \\ \gamma \end{pmatrix} = \mathbf{A}'\boldsymbol{\omega}$$

And we can simplify the previous expression of $f(\mathbf{w}, \gamma) = f(\boldsymbol{\omega})$ as follows:

$$f(\boldsymbol{\omega}) = \frac{1}{2} (\mathbf{A}'\boldsymbol{\omega} - \mathbf{y})^T (\mathbf{A}'\boldsymbol{\omega} - \mathbf{y})$$

In *Python*:

```
def loss_func (x):
    w = np.asarray(x).T

    cost = (1/2) * ((A @ w) - y).T @
            ((A @ w) - y)
    return np.squeeze(cost)
```

The gradient $\nabla f(\mathbf{w}, \gamma)$

Using the original expression, we have two variables, so we have to take derivatives from $f(\mathbf{w}, \gamma)$ with respect to both \mathbf{w}, γ . Arranging:

$$\begin{aligned} f(\mathbf{w}, \gamma) &= \frac{1}{2} (\mathbf{A}\mathbf{w} + \gamma\mathbf{e} - \mathbf{y})^T (\mathbf{A}\mathbf{w} + \gamma\mathbf{e} - \mathbf{y}) \\ &= \frac{1}{2} (\mathbf{w}^T \mathbf{A}^T \mathbf{A} \mathbf{w} + \gamma \mathbf{w}^T \mathbf{A}^T \mathbf{e} - \mathbf{w}^T \mathbf{A}^T \mathbf{y} \\ &\quad + \gamma \mathbf{e}^T \mathbf{A} \mathbf{w} + \gamma^2 \mathbf{e}^T \mathbf{e} - \gamma \mathbf{e}^T \mathbf{y} \\ &\quad - \mathbf{y}^T \mathbf{A} \mathbf{w} - \gamma \mathbf{y}^T \mathbf{e} + \mathbf{y}^T \mathbf{y}) \end{aligned}$$

$$\begin{aligned} \frac{\partial f(\mathbf{w}, \gamma)}{\partial \mathbf{w}} &= \mathbf{A}^T \mathbf{A} \mathbf{w} + \gamma \mathbf{A}^T \mathbf{e} - \mathbf{A}^T \mathbf{y} \\ &= \mathbf{A}^T (\mathbf{A} \mathbf{w} + \gamma \mathbf{e} - \mathbf{y}) \\ \frac{\partial f(\mathbf{w}, \gamma)}{\partial \gamma} &= \mathbf{e}^T \mathbf{A} \mathbf{w} + \gamma \mathbf{e}^T \mathbf{e} - \mathbf{e}^T \mathbf{y} \\ &= \mathbf{e}^T (\mathbf{A} \mathbf{w} + \gamma \mathbf{e} - \mathbf{y}) \end{aligned}$$

An easier expression of the gradient can be obtained using the simplified version $f(\boldsymbol{\omega})$:

$$\frac{\partial f(\boldsymbol{\omega})}{\partial \boldsymbol{\omega}} = \mathbf{A}'^T \mathbf{A}' \boldsymbol{\omega} - \mathbf{A}'^T \mathbf{y} = \mathbf{A}'^T (\mathbf{A}' \boldsymbol{\omega} - \mathbf{y})$$

We define it in *Python* as follows:

```
def loss_grad (x):
    w = np.asarray(x).T

    return A.T @ (A @ w - y)
```

The hessian $\nabla^2 f(\mathbf{w}, \gamma)$

Using the original expression we get a more complex expression:

$$H = \begin{pmatrix} \frac{\partial^2 f(\mathbf{w}, \gamma)}{\partial \mathbf{w}^2} & \frac{\partial^2 f(\mathbf{w}, \gamma)}{\partial \mathbf{w} \gamma} \\ \frac{\partial^2 f(\mathbf{w}, \gamma)}{\partial \gamma \mathbf{w}} & \frac{\partial^2 f(\mathbf{w}, \gamma)}{\partial \gamma^2} \end{pmatrix}$$

$$H = \begin{pmatrix} (\mathbf{A}^T \mathbf{A})_{p \times p} & (\mathbf{A}^T \mathbf{e})_{p \times 1} \\ (\mathbf{e}^T \mathbf{A})_{1 \times p} & (\mathbf{e}^T \mathbf{e})_{1 \times 1} \end{pmatrix}$$

We obtain a matrix by blocks. Using the simplified version of $f(\mathbf{w}, \gamma) = f(\boldsymbol{\omega})$ we obtain a much nicer expression:

$$\frac{\partial^2 f(\boldsymbol{\omega})}{\partial \boldsymbol{\omega}^2} = (\mathbf{A}'^T \mathbf{A}')_{(p+1) \times (p+1)}$$

In *Python* it is as simple as:

```
def loss_hess (x):
    w = np.asarray(x).T

    return A.T @ A
```

The inequality constraint $h(\mathbf{w})$

The *squared Euclidean norm* can be expressed as:

$$\|\mathbf{w}\|_2^2 = w_1^2 + w_2^2 + \dots + w_p^2 = \sum_{k=1}^p w_k^2$$

Since we are using $\boldsymbol{\omega} = \begin{pmatrix} \mathbf{w} \\ \gamma \end{pmatrix}_{(p+1) \times 1}$, we are going to ignore the last component setting it to zero:

$$\|\boldsymbol{\omega}\|_2^2|_{\gamma=0} = \|\mathbf{w}\|_2^2$$

So, in *Python*:

```
def cons_func (x):
    w = np.asarray(x).T
    w = w[:-1]
    return np.sum(w**2)
```

The Jacobian $\nabla h(\mathbf{w})$

We are going to use the same trick; we are going to set $\gamma = 0$. Thus, we get:

$$\frac{\partial}{\partial \omega_i} \sum_{k=1}^{p+1} \omega_k^2 = 2\omega = [2\omega_1, \dots, 2\omega_p, 0]$$

In *Python*:

```
def cons_jacobian (x):
    x[-1] = 0
    w = np.asarray(x).T

    return (2*w).tolist()
```

The hessian of the constraint $\nabla^2 h(\mathbf{w})$

We obtain a diagonal matrix with 2's on the diagonal except on the position $[(p+1), (p+1)]$ because of γ being 0, it is zero:

$$\left. \frac{\partial^2 \|\mathbf{w}\|_2^2}{\partial \mathbf{w}^2} \right|_{\gamma=0} = \begin{pmatrix} 2 & 0 & \dots & 0 \\ 0 & 2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 \end{pmatrix}$$

The way the minimization algorithm is implemented, it does not need a function that returns the raw hessian of the constraint(s), it needs a linear combination of the Hessians. Since we just have one constraint, and therefore, just one hessian, we should return the hessian multiplied by an arbitrary value v :

```
def cons_hess (x, v):
    x[-1] = 0
    w = np.asarray(x).T

    H = np.zeros((w.shape[0], w.shape[0]))
    np.fill_diagonal(H,2)
    H[-1,-1] = 0
    return v * H
```

Minimization: *death rate* data set

Now that all the ingredients are set up, we need to plug them into the minimization function from the *scipy.optimize* library:

```
# Define hyperparameter t:
t = 1

## OPTIMIZATION ##
# Assign initial weights:
w = np.ones(16)
# Define non linear constraint:
nonlinear_constraint =
    NonlinearConstraint(cons_func, -np.inf, t,
        jac = cons_jacobian, hess = cons_hess)
# Solve:
sol = minimize(loss_func, w, method='trust-
    constr', jac=loss_grad, hess=loss_hess,
    constraints=[nonlinear_constraint],
    options={'verbose': 1})
# Output:
print(sol)
```

Notes:

- We must define t beforehand, since it is a hyperparameter.
- We used a vector with zero components as the initial weights vector \mathbf{w} , since random initial weights with a seed did not lead to a good solution.

The complete output, using the *death rate* dataset, can be examined at the end of this paper (*Output 1*). Here, we are just going to show the found solution and the gradient:

```
# Solution (w):
([ 0.4882, -0.0464,  0.1029, -0.0377,
    0.0049, -0.0339, -0.2557,  0.0056,
    0.6497, -0.1262,  0.2134, -0.2078,
    0.1098,  0.3766,  0.0099, 895.141])
# gamma is the last component!
# Gradient norm: 0.21534461815763473
```

Note that minimizing without the constraint, we obtain the coefficients of the least squares linear regression (*unconstrained*), (*Output 2*):

```
# Solution (w):
([ 2.072, -2.177, -2.8319, -14.033,
   -115.329, -24.24, -1.145,  0.01,
    3.533,  0.5232,  0.268, -0.8889,
    1.866, -0.034,  0.533, 1862.39])
# gamma is the last component!
# Gradient norm: 0.08729604430708358
```

Conclusions

We have been able to implement the constrained ridge model in *Python* and solve successfully the optimization problem with the *death rate* data set.

References

J. CASTRO & F. J. HEREDIA, *Mathematical Optimization lecture notes*.

Ridge regression

Output 1

```
`xtol` termination condition is satisfied.
Number of iterations: 144, function evaluations: 191, CG iterations: 794, optimality: 3.97e-05, constraint violation: 0.00e+00, execution time: 0.27 s.
barrier_parameter: 2.0480000000000001e-09
barrier_tolerance: 2.0480000000000001e-09
cg_niter: 794
cg_stop_cond: 4
constr: [array([1.])]
constr_nfev: [191]
constr_nhev: [91]
constr_njev: [86]
constr_penalty: 1.0
constr_violation: 0.0
execution_time: 0.2699158191680908
fun: 61484.65464046687
grad: array([-1.05314313e+04, 1.00060115e+03, -2.22125180e+03, 8.13329504e+02,
-1.05852785e+02, 7.31601887e+02, 5.51735331e+03, -1.21838297e+02,
-1.40143362e+04, 2.72259783e+03, -4.60324321e+03, 4.48296568e+03,
-2.36932573e+03, -8.12422725e+03, -2.14833708e+02, 1.38982657e-05])
jac: [array([[ 0.9764791, -0.09277619, 0.20595548, -0.07541228, 0.00981472,
-0.06783446, -0.51157151, 0.0112969, 1.29941563, -0.25244051,
0.42681481, -0.41566262, 0.21968496, 0.75328205, 0.01991948,
0. ]])]
lagrangian_grad: array([-7.95691449e-06, -7.13667055e-06, 2.56749299e-05, -3.02727517e-06,
2.19963935e-05, 1.20500522e-06, -3.96506430e-05, -2.06478006e-06,
-1.83694065e-05, 8.83560779e-06, 1.73796743e-05, -1.15212924e-05,
-4.01049829e-06, -5.80644701e-06, 1.72957752e-05, 1.38982657e-05])
message: '`xtol` termination condition is satisfied.'
method: 'tr_interior_point'
nfev: 191
nhev: 86
nit: 144
niter: 144
njev: 86
optimality: 3.965064297517529e-05
status: 2
success: True
tr_radius: 3.0038937028194658e-09
v: [array([10785.1066819])]
x: array([ 4.88239551e-01, -4.63880973e-02, 1.02977740e-01, -3.77061410e-02,
4.90736021e-03, -3.39172299e-02, -2.55785757e-01, 5.64845110e-03,
6.49707815e-01, -1.26220254e-01, 2.13407404e-01, -2.07831310e-01,
1.09842480e-01, 3.76641024e-01, 9.95974041e-03, 8.95141404e+02])
```

Output 2

```
`xtol` termination condition is satisfied.
Number of iterations: 437, function evaluations: 426, CG iterations: 6106, optimality: 8.34e-02, constraint violation: 0.00e+00, execution time: 2.3 s.
cg_niter: 6106
cg_stop_cond: 4
constr: []
constr_nfev: []
constr_nhev: []
constr_njev: []
constr_penalty: 1.0
constr_violation: 0
execution_time: 2.34763503074646
fun: 23000.38378504645
grad: array([ 0.00508902, 0.000236, -0.00427007, -0.02073108, 0.00183165,
0.00412808, 0.00030917, 0.08338374, -0.00225027, 0.01116387,
-0.00093402, -0.00107135, 0.00372709, 0.00236066, 0.00010012,
-0.00470158])
jac: []
lagrangian_grad: array([ 0.00508902, 0.000236, -0.00427007, -0.02073108, 0.00183165,
0.00412808, 0.00030917, 0.08338374, -0.00225027, 0.01116387,
-0.00093402, -0.00107135, 0.00372709, 0.00236066, 0.00010012,
-0.00470158])
message: '`xtol` termination condition is satisfied.'
method: 'equality_constrained_sqp'
nfev: 426
nhev: 425
nit: 437
niter: 437
njev: 425
optimality: 0.0833837449317798
status: 2
success: True
tr_radius: 3.499966067021263e-09
v: []
x: array([ 2.07211884e+00, -2.17740576e+00, -2.83192146e+00, -1.40338963e+01,
-1.15329889e+02, -2.42401483e+01, -1.14517625e+00, 1.00428382e-02,
3.53300851e+00, 5.23275992e-01, 2.68007969e-01, -8.88934149e-01,
1.86632165e+00, -3.44450563e-02, 5.33898531e-01, 1.86239381e+03])
```